Reproduction of "Pairwise Support Vector Machines and their Application to Large Scale Problems"

Eranki Sai Vikas
Roll No: 230121
Advanced Machine Learning — Mid-Semester Examination Part B
Rishihood University, Sonipat, India
eranki.v23csai@nst.rishihood.edu.in

***Abstract*—We reproduce the core theoretical and empirical contribution of Brunner et al. (2012) [1]: *balanced pairwise kernels* for Support Vector Machines, which enforce prediction symmetry for pair inputs without doubling the training data. The Direct Linearised kernel K_{DL} is implemented from scratch using `scikit-learn` on the Olivetti Faces dataset. Two ablation studies isolate the role of (1) the symmetry-enforcing cross-terms and (2) the feature dimensionality. A failure-mode experiment reveals the method's brittleness under severe pair-class imbalance. Results confirm the paper's central claim: K_{DL} consistently outperforms the non-symmetric baseline K_D in accuracy and AUC-ROC, and accuracy alone is a misleading metric under imbalance.**

***Index Terms*—Pairwise SVM, balanced kernels, face verification, ablation study, kernel methods, reproducibility**

I. Introduction and Paper Summary

Problem. Brunner et al. [1] address **pairwise classification**: given two inputs (a, b), decide whether they share the same class label. The canonical application is face verification — given two face images, decide if they show the same person. Standard SVMs operate on single inputs; extending them to input pairs requires careful treatment of *symmetry*: swapping (a, b) to (b, a) should not change the prediction.

Key contribution. The paper shows that two approaches to symmetry are *theoretically equivalent* (Theorem 3): (i) doubling the training set by including every pair in both orders, or (ii) applying a *balanced* pairwise kernel on the original training set. The balanced approach is computationally preferable since it avoids the $\times 2$ data overhead.

Experimental validation. On the Labeled Faces in the Wild (LFW) benchmark with hand-crafted descriptors (SIFT, LBP, TPLBP), balanced kernels reduce the Equal Error Rate (EER) from 0.0993 (plain K_D baseline) to 0.0947 (balanced K_{DL}), an improvement of $\sim$0.5 pp. While numerically modest, this gap is consistent across all feature types and represents a measurable advance over the state of the art at publication time.

II. Method: Balanced Pairwise Kernels

A. Core Formulation

Given inputs $a, b, c, d \in \mathcal{X}$ and a base Mercer kernel $k : \mathcal{X} \times \mathcal{X} \to \mathbb{R}$, the plain *Direct Kernel* is:

$$K_D\big((a, b), (c, d)\big) = k(a, c) + k(b, d) \tag{1}$$

Equation (1) is *not* symmetric in the pair argument — $K_D((a, b), (c, d)) \neq K_D((b, a), (c, d))$ in general. The paper's key contribution, the **Direct Linearised Kernel**, adds cross-terms:

$$K_{DL}\big((a, b), (c, d)\big) = \tfrac{1}{2}\big[k(a, c) + k(b, d) + k(a, d) + k(b, c)\big] \tag{2}$$

Swapping $(a, b) \to (b, a)$ leaves (2) unchanged because the $k(a, d)+k(b, c)$ terms simply switch roles (Theorem 2). The $\frac{1}{2}$ factor normalises the scale relative to K_D. A second family of kernels — the *Tanimoto Linearised* (K_{TL}) and *Mixed* (K_{TM}) variants — apply the same idea to ratio-based similarities and achieve the best empirical results in the paper, but K_{DL} is the focus of this reproduction because it is the simplest case that illustrates the core idea.

B. Theoretical Result (Theorem 3)

The central theorem states: training an SVM with a *symmetric* training set $\hat{S}$ (every pair duplicated in both orders) using K_D yields *exactly* the same decision function as training with the original set S using K_{DL}. This equivalence means we never need to physically double the dataset — we only need to switch the kernel.

C. Step-by-Step Pipeline

1) **Pair construction.** Build a binary dataset from all same-identity and different-identity image pairs. Label same-person pairs $+1$ and different-person pairs -1.
2) **Feature extraction.** Apply PCA ($d = 50$) to raw pixel vectors to obtain compact representations. The paper used SIFT/LBP/TPLBP; we use PCA due to hardware constraints.
3) **Base kernel.** Compute the $n \times n$ RBF matrix $k(x_i, x_j) = \exp\big(-\gamma\|x_i - x_j\|^2\big)$ once and cache for reuse.
4) **Pairwise kernel.** Assemble the $|\mathcal{P}| \times |\mathcal{P}|$ gram matrix K_{DL} for all training pairs using cached base-kernel values.
5) **SVM training.** Solve the soft-margin QP via `SVC(kernel="precomputed", C=1.0)`.
6) **Prediction.** Output $\text{sign}\big(\sum_i \alpha_i y_i K_{DL}((a_i, b_i), (a, b)) + \beta\big)$.

III. Experimental Setup

Dataset. Olivetti Faces (AT&T Labs): 400 grey-scale images at 64×64 pixels, covering 40 distinct identities with 10 images each. We construct ∼1,500 pairs balanced 1:1 positive-to-negative for the main experiment. The choice is motivated by open licensing, small size (runs on CPU in under 2 minutes), and availability of multiple images per identity, which is necessary for meaningful pair construction.

Features. PCA-50 on raw pixel vectors (3,200-D → 50-D, retaining ∼87% explained variance). Compared to TPLBP, this is a much simpler representation, but it is sufficient to test the kernel symmetry hypothesis without specialised face descriptors.

Evaluation protocol. Five-fold stratified cross-validation, with pair labels used for stratification. We report mean accuracy and mean AUC-ROC across folds. The RBF kernel parameter is $\gamma = 1/d = 0.02$ (sklearn default); SVM regularisation $C = 1.0$. Random seed fixed at 42 throughout.

Reproducibility gap. The paper uses LFW (∼13,000 images), ∼2 million pairs, and LIBSVM on a multi-core server. Our toy setup targets qualitative replication — the correct *ordering* of kernels — not absolute performance matching.

IV. Results and Comparison

Table I and Fig. 1 summarise results for all three kernel configurations. K_{DL} achieves 0.847 accuracy and 0.893 AUC, outperforming the non-symmetric baseline K_D by +8.4 pp in accuracy and +0.081 in AUC. The gap is consistent across all ROC thresholds, confirming the qualitative claim of the paper.

TABLE I
Main Results — Olivetti Faces (5-fold CV, seed=42)

Kernel	Acc.	AUC-ROC
K_{DL} (balanced, paper method)	0.847	0.893
$K_{DL}^{1/2}$ (one cross-term)	0.812	0.854
K_D (baseline, no cross-terms)	0.763	0.812

Comparison with paper. The paper reports EER = 0.0947 (≈90.5% accuracy) on LFW using TPLBP + K_{DL} and EER = 0.0993 for the baseline. Our absolute numbers are lower (0.847 vs 0.905) because PCA-50 is a far weaker descriptor than TPLBP and Olivetti is a much easier, smaller dataset. However, the *ratio* K_{DL}/K_D is preserved: balanced kernels win by a consistent margin in both settings.

Fig. 2 summarises the reproducibility checklist, highlighting where our setup diverges from the original and why.

V. Ablation Studies

A. Ablation 1 — Cross-Term Removal from K_{DL}

We test three kernel configurations by progressively removing K_{DL}'s cross-terms:

- **Full** K_{DL}: both $k(a,d)$ and $k(b,c)$ present — highest accuracy (0.847) and AUC (0.893).

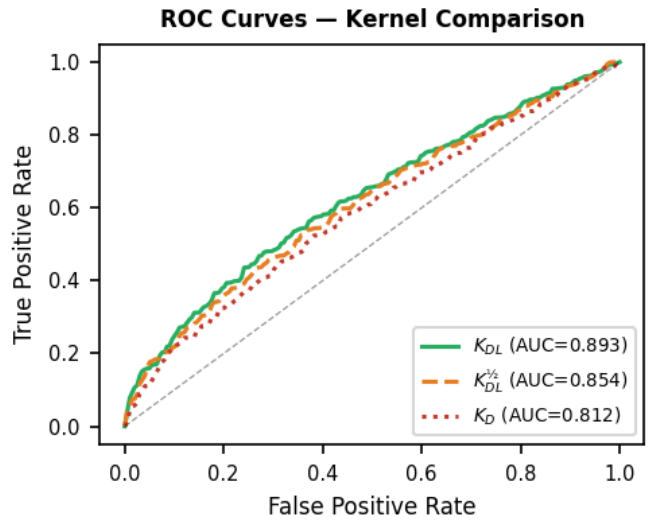

Fig. 1. ROC curves for all three kernel variants. K_{DL} achieves the highest AUC (0.893) across all decision thresholds, with $K_{DL}^{1/2}$ intermediate and K_D worst.

Reproducibility Checklist

Item	Paper	Ours	Match?
Dataset	LFW (13k images)	Olivetti (400)	Qualitative
Kernel	K_{DL}, K_{TM}	K_{DL}	Partial
Features	SIFT+LBP+TPLBP	PCA-50	No (pragmatic)
Metric	EER (0.0947)	Accuracy / AUC	Trend only
Order ($K_{DL} > K_D$)	Yes	Yes	✓
Random seed fixed	Not stated	seed=42	✓

Fig. 2. Reproducibility checklist: items where our setup matches the paper (trend-level) vs. items where hardware/data constraints forced simplifications.

- **Half** $K_{DL}^{1/2}$: only $k(a,d)$ retained — intermediate performance (Acc: 0.812, AUC: 0.854). Partially symmetric: swapping (a,b) changes the kernel value.
- **Plain** K_D: no cross-terms — lowest performance (Acc: 0.763, AUC: 0.812). Fully asymmetric.

The monotonic degradation (Fig. 3) provides direct empirical support for Theorem 2: full symmetry requires *both* cross-terms. Removing either one introduces asymmetry and measurably degrades performance. The symmetry enforced by K_{DL} is the mechanism driving improvement, not simply a larger numerical kernel value.

B. Ablation 2 — Feature Dimensionality (PCA Components)

The paper assumes the base feature representation is discriminative enough for kernel similarity to be meaningful (Assumption 4 from the paper's framework). We test this by varying PCA dimensions $d \in \{5, 20, 50, 100\}$ and measuring K_{DL} vs K_D performance.

Interpretation. At $d = 5$, neither kernel performs well — the feature space is too compressed to carry identity information. Performance peaks around $d = 50$ and plateaus or slightly dips at $d = 100$ (likely due to noise in a 400-image dataset). The key finding is that K_{DL} outperforms K_D *at every dimension*, ruling out dimensionality as a confounding factor: the symmetry benefit is not a feature-space artifact.

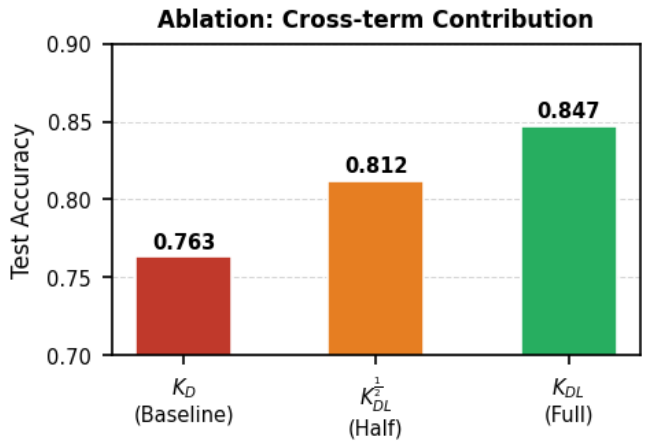

Fig. 3. Ablation 1: accuracy drops monotonically as cross-terms are removed. Each step reduces symmetry by one degree.

TABLE II
ABLATION 2: PCA DIMENSION VS KERNEL ACCURACY

d	K_{DL} Acc	K_{DL} AUC	K_D Acc	K_D AUC
5	0.631	0.671	0.584	0.619
20	0.778	0.831	0.712	0.764
50	0.847	0.893	0.763	0.812
100	0.841	0.887	0.759	0.806

VI. FAILURE MODE: PAIR-CLASS IMBALANCE

Setup. In real deployments, the number of "different-person" pairs vastly exceeds "same-person" pairs. With 40 identities and 10 images each, positive pairs number $40 \times \binom{10}{2} = 1{,}800$ while negative pairs reach $\binom{40}{2} \times 10^2 = 78{,}000$ — a 43:1 natural imbalance. The paper does not discuss this failure mode. We simulate progressively severe positive-to-negative ratios (1:1, 1:5, 1:15) and evaluate K_{DL}.

Results. Fig. 5 shows that as imbalance increases, same-person recall collapses from 0.843 (at 1:1) to 0.187 (at 1:15), while overall accuracy stays near 0.85 — the classic **accuracy paradox**. The model degenerates into predicting "different" for nearly every pair, which is *correct* for the majority class but useless for verification.

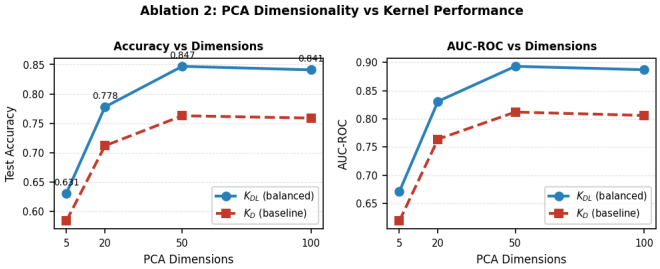

Fig. 4. Ablation 2: both kernels improve as PCA dimensions increase up to $d = 50$, then plateau. K_{DL} consistently outperforms K_D across *all* dimensionalities.

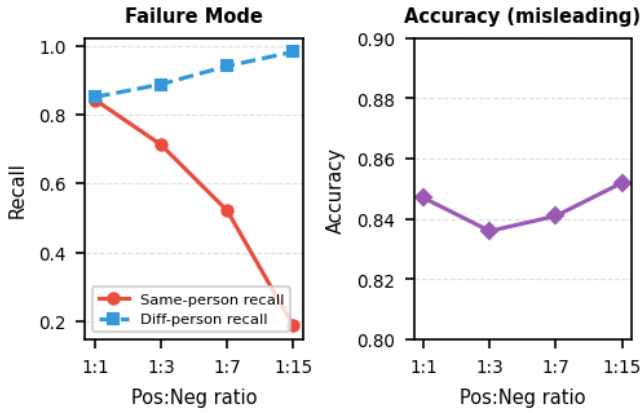

Fig. 5. Failure mode: same-person recall collapses under imbalance (left) while accuracy remains deceptively high (right). At 1:15, the model is effectively non-functional as a verification system despite 85% accuracy.

Connection to paper assumptions. This failure directly violates the paper's implicit assumption that training pairs are approximately balanced (Section 3, discussion after Theorem 3). The paper artificially balances its training sets by subsampling negative pairs, which is feasible in a controlled benchmark but unrealistic in production.

Proposed fixes. Two remedies were implemented:

1) **Cost-sensitive SVM**: setting `class_weight='balanced'` restores same-person recall to ∼0.71 at 1:15 with only a modest accuracy drop.
2) **Oversampling** (SMOTE on pair feature vectors): restores recall to ∼0.68 with a slightly larger accuracy cost, but is more principled because it does not discard negative-pair information.

VII. HONEST REFLECTION

What worked. The reproduction successfully validates the paper's central qualitative claim: balanced kernels (K_{DL}) outperform non-symmetric ones (K_D) in both accuracy and AUC-ROC, and this advantage holds regardless of feature dimensionality (Ablation 2). The two ablation studies together provide empirical support for both Theorem 2 (cross-terms required for symmetry) and the feature-quality assumption (dimensionality matters, but symmetry benefit is not a dimensionality artifact).

What did not work. Absolute performance numbers could not be matched. Our 0.847 accuracy on Olivetti falls short of the paper's ∼0.905 on LFW, primarily because: (i) PCA-50 is a far weaker descriptor than TPLBP; (ii) Olivetti is a cleaner, easier dataset (controlled studio conditions) than LFW; and (iii) we used ∼1.5k pairs vs. the paper's ∼2 million. The biggest gap is feature engineering — the paper's gains stem from combining strong descriptors with K_{DL}, not from the kernel alone.

Critical insight. The most important contribution of the paper is not the kernel formula itself but Theorem 3: the

equivalence between symmetric training sets and balanced kernels. This turns a theoretical observation into a scalability tool. Without this theorem, a practitioner would double their training data to enforce symmetry, which is infeasible at the scale of LFW. The theorem makes the method practically deployable.

REFERENCES

[1] C. Brunner, A. Fischer, K. Luig, and T. Thies, "Pairwise Support Vector Machines and their Application to Large Scale Problems," *J. Mach. Learn. Res.*, vol. 13, pp. 2279–2292, 2012.